



InDeF: An Advanced Defragmenter Supporting Migration Offloading on ZNS SSD

Wenjie Qi, Zhipeng Tan, Jicheng Shao, Lihua Yang, Yang Xiao

*Wuhan National Laboratory for Optoelectronics,
Huazhong University of Science and Technology (HUST), China*





Outline

- **Background and motivation**
- **Our Work: InDeF**
- **Performance Evaluation**
- **Conclusion**

Background : What is ZNS SSD?

- The logical address space is divided into fixed-sized zones
- Each zone must be written sequentially and reset explicitly for reuse




ZNS: Avoiding the Block Interface Tax for Flash-based SSDs

Matias Björling*, Abutalib Aghayev^o, Hans Holmberg*, Aravind Ramesh*, Damien Le Moal*,
Gregory R. Ganger[‡], George Amvrosiadis[†]

*Western Digital ^oThe Pennsylvania State University [†]Carnegie Mellon University



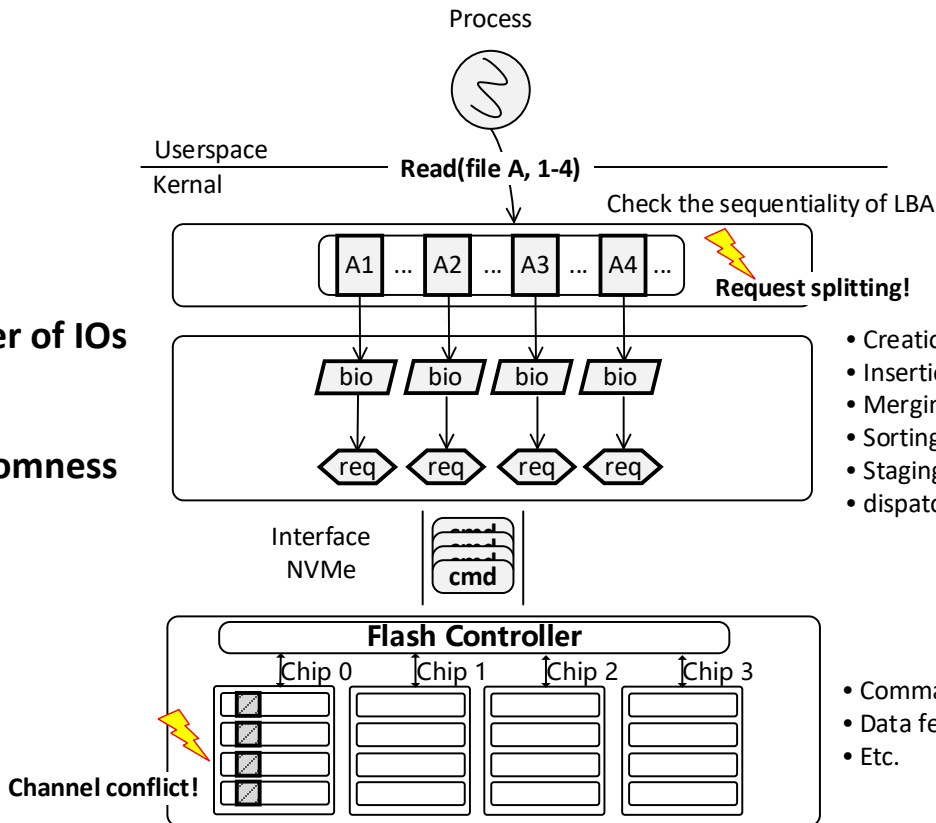
Zoned Namespace Command Set Specification



Background : What is fragmentation?

Request splitting

- 1) Increases the number of IOs
- 2) Makes I/Os smaller
- 3) Increases their randomness



```
if (bio && !page_is_mergeable(F2FS_I_SB(inode), bio,
    *last_block_in_bio, block_nr)) {
    submit_and_realloc:
    __submit_bio(F2FS_I_SB(inode), bio, DATA);
    bio = NULL;
}
```

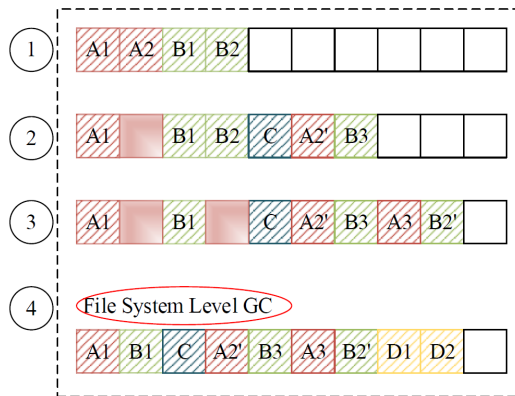
Logical Fragmentation

- Creation
- Insertion
- Merging
- Sorting
- Staging
- dispatching

Physical Fragmentation


- Command processing
- Data fetching(DMA)
- Etc.

Motivation : Fragmentation accumulation



➤ Incoming Data Stream :

➤ A1A2, B1B2 → C, A2', B3 → A3, B2' → D1D2



➤ Fragmentation on ZNS SSD

Sequential read after running Fileserver/Varmail
Filesystem: F2FS

Motivation : Definition of the fragmentation

➤ Definition of the fragmentation of an I/O request

➤ **The degree of logical fragmentation (DoLF)** is the number of logical fragments in an I/O range

➤ **The degree of physical fragmentation (DoPF)**

We measure the degree of physical fragmentation of an I/O request by how evenly the data in the I/O range are distributed among the flash parallel units.

$$DoPF = \frac{\sum_{i=1}^L \left(N_i - \frac{M}{L}\right)^2}{L}$$

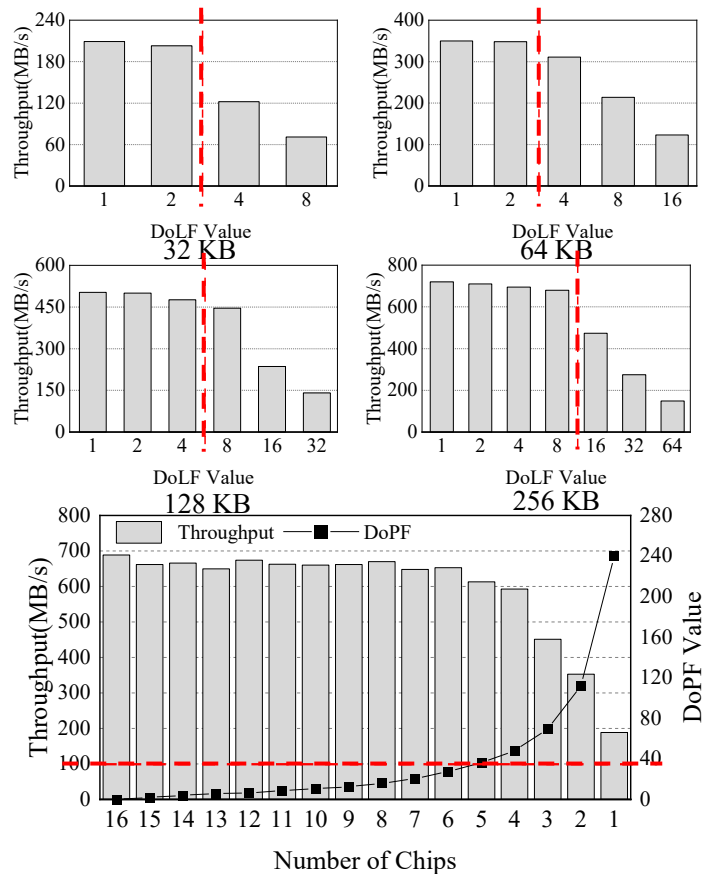
Motivation : The impact of fragmentation

➤ Evaluation Setup

- 32KB 64KB 128KB 256KB O_DIRECT sequential read on F2FS
- Varying DoLF/DoPF

➤ Observations

- A low DoLF value has a small impact on I/O performance
 - Low kernel overhead
- When the DoPF value is small, e.g., less than 40, physical fragmentation has a negligible impact on I/O performance
 - Software overhead of I/O dominates the total I/O latency



No need to defragment all the fragments!

Motivation : How to Select the Appropriate Data

- Modern storage systems typically perform data access in non-uniform distribution [1],[2]
- Multiple reads to fragmented data accumulate the access latency caused by fragmentation
- **Our idea:**
 - Defragmenting data with a low degree of fragmentation or cold data that is rarely accessed provides little performance gain
 - **We define the I/O data defragmentation priority (IODP)**

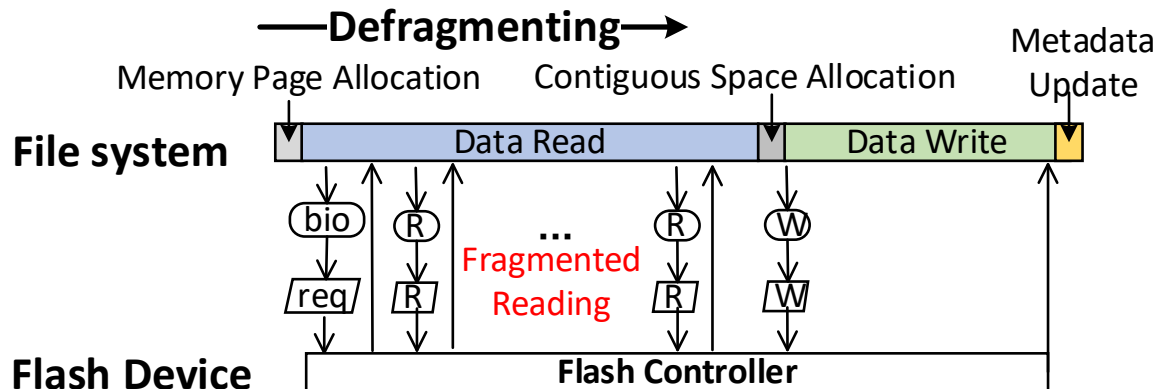
$$IODP = (\alpha \cdot DoLF + \beta \cdot DoPF) \times readcount$$

[1] Q. Wang, J. Li, P. P. Lee, T. Ouyang, C. Shi, and L. Huang, “Separating data via block invalidation time inference for write amplification reduction in log-structured storage,” in *Proc. of USENIX FAST*, 2022.

[2] Y. Lv, L. Shi et al., “Access characteristic guided partition for read performance improvement on solid state drives,” in 2020 57th ACM/IEEE Design Automation Conference (DAC). IEEE, 2020, pp. 1–6.

Motivation : The Conventional Defragmenter

- Cause a significant increase in the **host memory usage** and invoke page frame reclamation
- Result in a large **chip idle interval** in the SSD
- Migrate **the entire contents of files** even when there are few fragments
- The additional writes **reduce the lifespan** of modern storage devices
- **Degrades the performance** of co-running applications
- Time-consuming

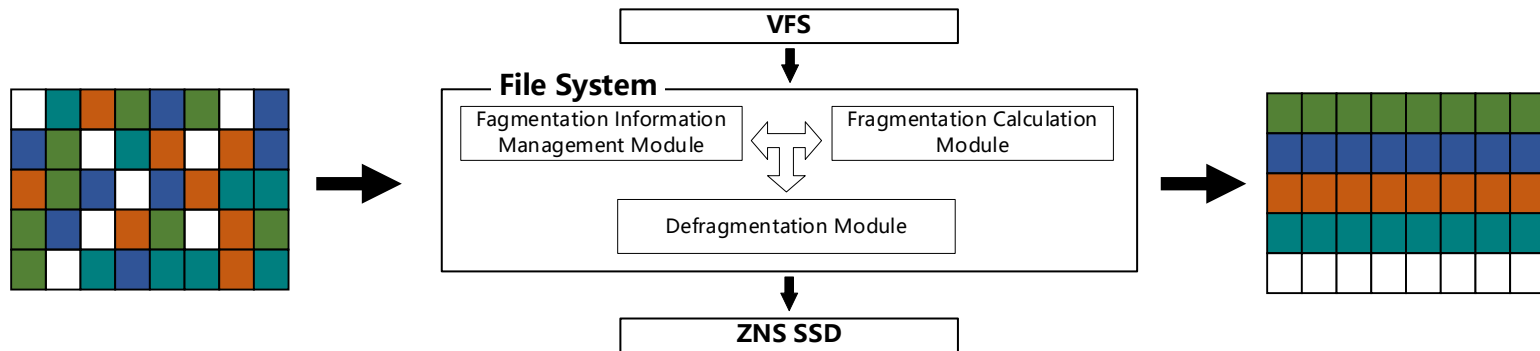


Our Scheme: InDeF for Zoned Namespace SSD

➤ Main goals

- ✓ Minimizes the amount of migration data for defragmentation to reduce the write traffic of the underlying device
- ✓ Decreases elapsed time of defragmentation to reduce the impact on co-running application performance

InDeF for Zoned Namespace SSD



➤ The Fragmentation Information Management Module

Collects I/O information from the filesystem

Manages I/O fragmentation information

➤ The Fragmentation Calculation Module

Calculates the DoLF value and DoLP value based on the collected I/O information

➤ The Defragmentation Module

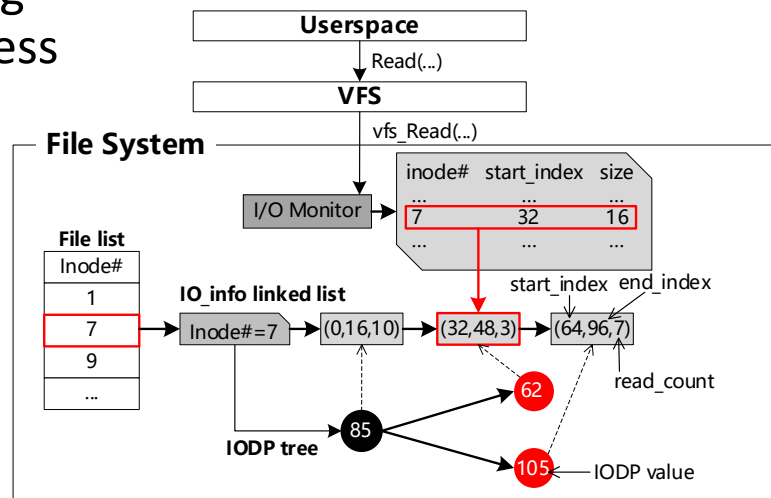
Filters the fragments based on the I/O fragmentation information

Offloads the data migration from the host to the SSD

InDeF for Zoned Namespace SSD

➤ The Fragmentation Information Management Module

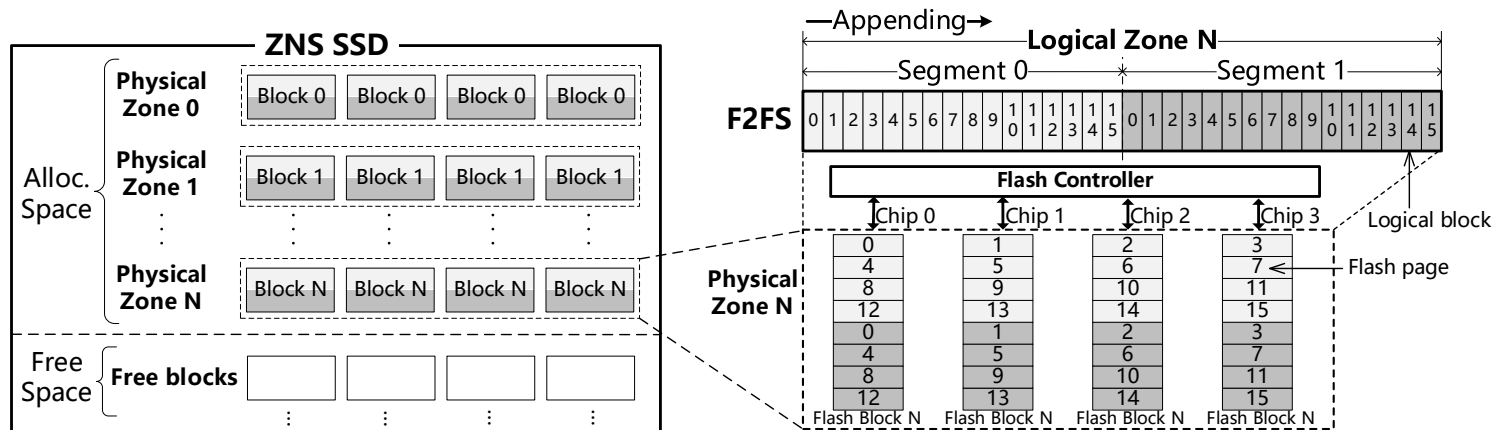
- **Monitors** I/O activity at the file system layer
- **Creates** file list and I/O information linked list
- **Merges** I/O requests that have overlapping addresses and preserves the largest address range
- **Inserts** the IODP of each I/O information into a red-black tree (called IODP tree) in order



InDeF for Zoned Namespace SSD

➤ The Fragmentation Calculation Module

- Gets the physical location of the logical block inside the device by the segment number and in-segment offset of the logical block
- The i -th logical block of a segment is stored on the j -th flash chip that $j = i \% (\text{the number of parallel flash chips})$



InDeF for Zoned Namespace SSD

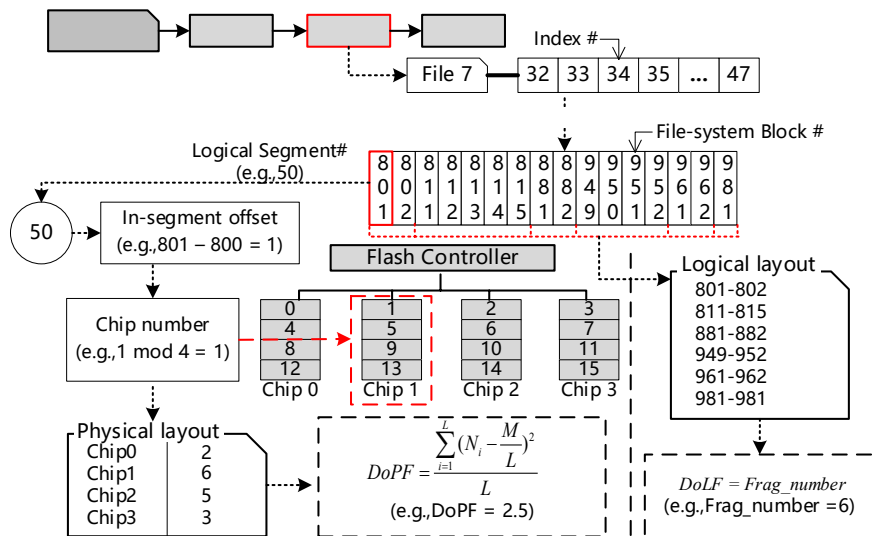
➤ The Fragmentation Calculation Module

➤ An example of calculating the DoLF and the DoPF

The DoLF can be calculated by the number of logical fragments in the I/O range

The DoPF can be calculated by

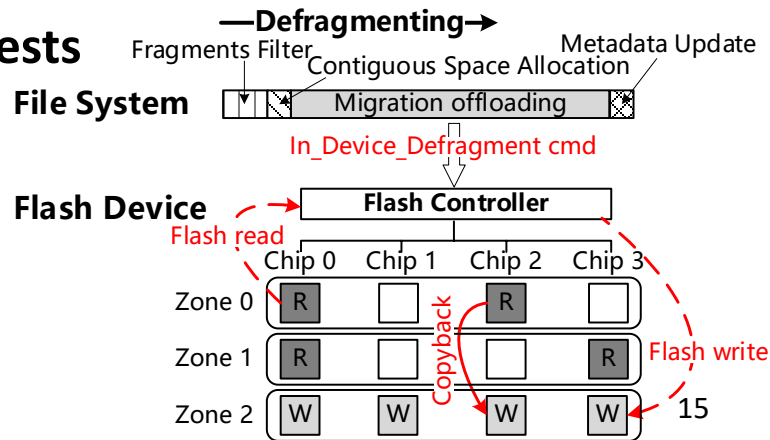
$$DoPF = \frac{\sum_{i=1}^L (N_i - \frac{M}{L})^2}{L}$$



InDeF for Zoned Namespace SSD

➤ The Defragmentation Module

- **Truncates** I/O information entries with a low IODP value based on the ordered IODP tree
- **Groups** data based on whether the data page is dirty or not
- **Sends** the In_Device_Defragment command
 - Contains a set of source LBAs and a set of destination LBAs
- **Handles** migration offloading and host requests
 - Flash read and write
 - Copyback



Performance Evaluation

➤ Experimental Setup

- InDeF emulator based on FEMU

The CASE of FEMU:

Cheap, Accurate, Scalable and Extensible Flash Emulator

Huaicheng Li, Mingzhe Hao, Michael Hao Tong,
Swaminathan Sundararaman[†], Matias Björling[‡], Haryadi S. Gunawi
University of Chicago [†]Parallel Machines [‡]CNEX Labs

- A QEMU-based and DRAM-backed NVMe SSD Emulator
- <https://github.com/ucare-uchicago/femu>

➤ Comparison

- defrag.f2fs[3] vs. FragPicker[4] vs. InDeF

➤ Workloads

- Synthetic and Macro Benchmarks

➤ Objectives

- Does InDeF reduce the amount of writes for defragmentation?
- Does InDeF achieve a similar level of performance gain, compared with conventional tools?
- Does InDeF decrease the elapsed time of defragmentation?

[3] “defrag.f2fs,” 2022, <https://manpages.debian.org/testing/f2fs-tools/defrag.f2fs.8.en.html>.

[4] J. Park and Y. I. Eom, “Fragpicker: A new defragmentation tool for modern storage devices,” in Proceedings of the ACM SIGOPS 28th Symposium on Operating Systems Principles, 2021, pp. 280–294.


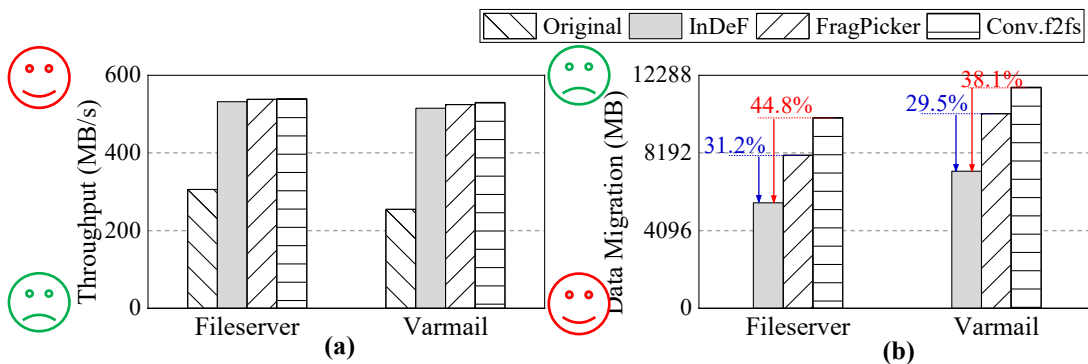
Experiment – Performance and Write Amount

➤ Performance

- InDeF improve the throughput by about 72% (filesystem) and 112% (varmail), compared with that before defragmentation

➤ Write amount


- InDeF reduces the amount of writes by around 31.2%-44.8%(Fileserver) and 29.5%-38.1%(Varmail)



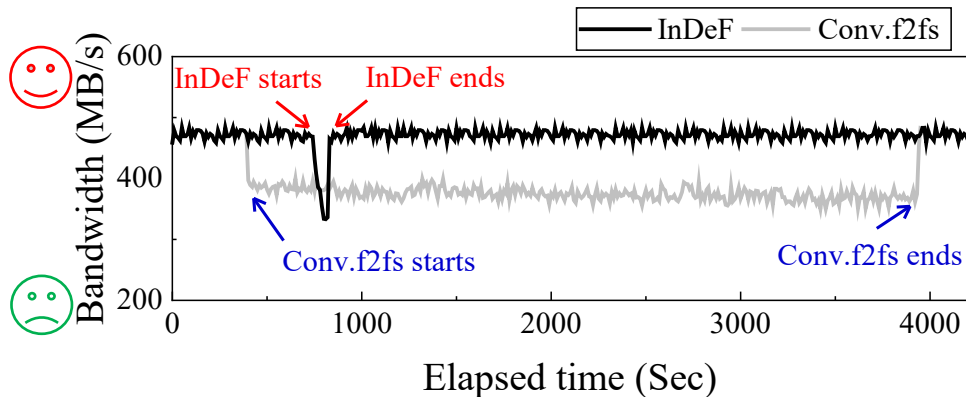
Experiment - Elapsed time of defragmentation

➤ Elapsed time of defragmentation

- InDeF increased chip utilization during migration by 8.9x (filesystem) and 7.4x (varmail)
- InDeF decreases the elapsed time of defragmentation time by 94.2% (filesystem) and 91.6% (varmail) due to the higher chip utilization



(a)



(b)

More experiments in our paper!

- **Synthetic Benchmarks**

- **Database workloads**

 - RocksDB YCSB-C

 - ...

Conclusion

➤ Main goals

- ✓ Minimizes the amount of migration data for defragmentation to reduce the write traffic of the underlying device
- ✓ Decreases elapsed time of defragmentation to reduce the impact on co-running application performance

➤ InDeF for ZNS SSD

- ✓ Combines the degree of fragmentation and access hotness to find out the most suitable data set to migrate in each file
- ✓ By the In_Device_Defragment command, InDeF offloads the data migration from the host to the ZNS SSD to improve the efficiency of defragmentation
- ✓ Decreases the elapsed time of defragmentation significantly while minimizing the amount of data migration for defragmentation

Thanks!

